RAHUL M

Data Scientist | Bengaluru, Karnataka

+91 8220250263 | rahul502mv@gmail.com | LinkedIn | GitHub | Portfolio

PROFESSIONAL SUMMARY

I'm an entry-level Data Scientist with a B.Tech in Information Technology, and I've spent the last year or so applying machine learning and statistical modeling to actual business problems — not just notebooks. Python, SQL, and predictive algorithms (Scikit-learn, LightGBM, XGBoost) are where I'm strongest, and I've also built an LLM-powered tool using the Groq API and Llama 3.3 70B for natural-language data querying. I've taken at least one project the full distance, from raw data to a deployed app, and I'm comfortable working across teams to turn complicated data into something people can actually act on.

TECHNICAL SKILLS

Programming & Algorithms: Python (Pandas, NumPy, Scikit-learn, Matplotlib, Seaborn, SHAP), machine learning algorithms for classification and risk prediction
Machine Learning & Statistical Modeling: Logistic Regression, Random Forest, XGBoost, LightGBM, Isolation Forest, statistical analysis, model evaluation
Applied LLM / GenAI: Groq API, Llama 3.3 70B, Natural Language-to-SQL, prompt-driven data workflows
Data Science: Data cleaning, EDA, feature engineering, feature scaling, predictive analytics
Database: SQL, Oracle (Joins, Subqueries, CTEs, Window Functions, Stored Procedures), SQL Server
Visualization & BI: Power BI (DAX, Power Query, Data Modeling, Star Schema), Tableau, Excel (Pivot Tables, VLOOKUP, XLOOKUP, Macros/VBA)
Deployment & Tools: FastAPI, Streamlit, Docker, Git/GitHub, Render, Jupyter/Colab

PROJECTS

Enterprise AI-Powered Credit Risk Intelligence Platform | Live Demo — *Python, Pandas, Scikit-learn, LightGBM, SHAP, FastAPI, Streamlit, Docker*

- Built an end-to-end ML pipeline on the Home Credit Default Risk dataset — 307,500+ records, 122+ features — using LightGBM and XGBoost to forecast loan default risk, landing at an ROC-AUC of 0.66.
- Added SHAP explainable AI on top so predictions translate into feature-level reasons a non-technical stakeholder can actually follow.
- Took it all the way to deployment: a FastAPI + Streamlit app running on Docker via Render, handling real-time predictions and CSV batch uploads.

AI SQL Query Generator — *Python, Groq API, Llama 3.3 70B, SQL, Streamlit*

- Built a Natural Language-to-SQL app on top of the Groq API and Llama 3.3 70B — type a business question, get back an executable SQL query.
- The pipeline covers generation, execution, and result analysis, so the AI-generated insight comes out the other end without manual steps in between.
- Streamlit ties it together as an interactive app, which turned out to be a nice, concrete way to show applied LLM integration in a data science context.

Netflix Content Analytics Dashboard — *Python, Pandas, NumPy, Matplotlib, Seaborn*

- Analyzed 8,800+ Netflix titles and built 10+ visualizations around genre, rating, and country-wise content trends.
- The cleaning and feature engineering work underneath is what made the EDA useful — it's what ended up informing the content-strategy recommendations.

Blinkit Sales Analysis Dashboard — *SQL, Power BI, Excel, Power Query*

- Built an interactive Power BI dashboard on 8,500+ sales records and found that Tier-3 outlets were quietly the top revenue contributor.
- Redesigned the SQL-driven data models and automated the dashboard refreshes, which cut reporting turnaround by 40%.

INTERNSHIP EXPERIENCE

Data Science Intern — QSpiders, Bengaluru Jan 2026 – Jul 2026

- Delivered real-world data science projects end to end using Python, SQL, and machine learning, working alongside a cross-functional training team the whole way.
- Built ML models and dashboards by running data preprocessing, EDA, feature engineering, and predictive modeling on applied business problems.

Data Analyst Intern — Gateway Software Solutions, Chennai Jul 2025 – Aug 2025

- Automated SQL and Power BI reporting across 6 KPIs — cut manual reporting time by 35%.
- EDA and statistical analysis on customer and sales data helped speed up decision-making by roughly 15%.
- Redesigned the star-schema data models in Power BI, which brought query load down by 20%.

EDUCATION

Bachelor of Technology, Information Technology
Shreenivasa Engineering College, Tamil Nadu | 2022 – 2026 | CGPA: 7.9 / 10.0

CERTIFICATIONS

- Data Science with Python — QSpiders (2026)
- Data Science for Engineers — NPTEL (2025)
- IBM AI Fundamentals — IBM SkillsBuild (2025)
- Soft Skill Development Certificate

CORE COMPETENCIES

Analytical & Strategic Thinking • Problem Solving • Cross-Functional Collaboration • Communication • Time Management